

Beaker Specification

Jacob Payne Jake O'Shannessy Alexey Troitskiy

March 13, 2019

Abstract

We describe a secure and extensible operating system for smart contracts. Using a capability-based exokernel protocol, we can facilitate secure isolation, perform upgrades in a secure and robust manner, and prevent privilege escalation at any point in the development process. The protocol is intended to serve as an open standard and building block upon which more advanced and tailored security models can be built.

Contents

1	Definitions	4
1.1	Kernel Instance	4
2	Kernel Storage	4
3	Procedures	4
3.1	Procedure Key	4
3.2	Key Index	5
3.3	Procedure List	5
3.4	Procedure Heap	6
4	System Calls and Capabilities	7
4.1	Executing a System Call	8
4.2	Return and Error Codes	8
4.3	Capability Subsets	9
4.4	Call System Call	9
4.4.1	System Call Format	9
4.4.2	Capability Format	10
4.4.3	Capability Subsets	10
4.4.4	Error Codes	11
4.5	Register Procedure System Call	11
4.5.1	System Call Format	11
4.5.2	Capability Format	12
4.5.3	Capability Subsets	12
4.5.4	Function	13
4.5.5	Error Codes	14
4.6	Delete Procedure System Call	14
4.6.1	System Call Format	14
4.6.2	Capability Format	14
4.6.3	Capability Subsets	15
4.7	Set Entry Procedure System Call	15
4.7.1	Error Codes	15
4.7.2	System Call Format	16

4.7.3	Capability Format	16
4.7.4	Capability Subsets	16
4.7.5	Function	16
4.7.6	Error Codes	17
4.8	Write System Call	17
4.8.1	System Call Format	17
4.8.2	Capability Format	17
4.8.3	Capability Subsets	18
4.8.4	Error Codes	18
4.9	Log System Call	18
4.9.1	System Call Format	18
4.9.2	Capability Format	19
4.9.3	Capability Subsets	19
4.9.4	Error Codes	20
4.10	Value Send System Call	20
4.10.1	System Call Format	20
4.10.2	Capability Format	20
4.10.3	Capability Subsets	20
5	Procedure Bytecode Validation	21

1 Definitions

1.1 Kernel Instance

A kernel instance is an Ethereum contract which using the Beaker kernel. It is this contract that holds the storage data etc. of the system.

2 Kernel Storage

This is a critical feature. *All* data stored by the kernel is stored in a region of *storage* called *kernel storage*. *No* other data can be stored in *kernel storage*.

This kernel storage is defined as all storage locations within a certain storage range. All values in kernel storage share the same storage key prefix. This prefix is currently 0x00, and so currently includes all values in the interval [0x0000...0000, 0x0100...0000) (some zeros have been elided for typographic reasons, but each values is 32 bytes).

Table 1: Overview of kernel storage layout.

0x	00	01 - 04	05 - 1f	1e - 20	Description
0x	00	00 00 00 00	Proc Key	...	Procedure Heap
0x	00	00 00 00 01	00 ... 00	00 00 00	# of Procedures
0x	00	00 00 00 01	Key Index	00 00 00	Procedure List
0x	00	00 00 00 02	00 ... 00	00 00 00	Kernel Address
0x	00	00 00 00 03	00 ... 00	00 00 00	Current Procedure
0x	00	00 00 00 04	00 ... 00	00 00 00	Entry Procedure

3 Procedures

Data about procedure is store in 2 locations: the procedure table, and the procedure heap. The procedure table is simply a managed, enumerable, and iterable list of procedure keys (procedure key being something that identifies a particular procedure). The data associated with a procedure, such as its Ethereum address and its capabilities are stored on the procedure heap. Given a procedure key, the data associated with a procedure can be located using that key.

3.1 Procedure Key

Each procedure is defined by a procedure key. This key is a sequence of 24 8-bit bytes. This is treated as a sequence of 192 bits in all cases.

The kernel itself treats the procedure keys as an opaque 192-bit value, however, the capability system and some functions apply prefixes to these key value. For example: capabilities that rely on procedure key values will often use a prefix value to define a range of keys. This allows one to defined a hierarchy of procedures, much like a directory structrue. See those capabilities for more detail.

3.2 Key Index

Each procedure key included in the kernel is associated with a key index, which identifies a procedure in the procedure table. This is not used anywhere except the procedure table itself. It is 1-based, that is the first value is 1, and the value 0 is a null value.

As shown below, the maximum number of procedures is $2^{24} - 1 = 16,777,215$, therefore the maxium value of the key index is 16,777,215, therefore the key index lies in the range $[1, 1677215]$.

For this reason the key index is specified as 24 bytes.

3.3 Procedure List

The procedure list is simply an array of storage values. The first value is the length of this list, and the subsequent values are the procedure keys of the list.

When a procedure is added to the kernel:

1. The procedure data is added to the procedure heap (see Section 3.4).
2. The procedure key is appended to the end of the array.
3. The length value (the first value) is incremented by one.

When a procedure is deleted from the kernel:

1. If the procedure key is the same as the Entry Procedure Key, abort and throw an error.
2. If the procedure key does not exist in the list, abort and throw an error.
3. The length value is decremented by one.
4. If the procedure being deleted is not the last in the list, the last in the list is copied to overwrite the key being deleted. This also accounts for the case of an empty list.

It is important to note that none of these steps consider deletion (zeroing) of data. This is optional for efficiency and can be performed any time after the length value has been decremented.

The procedure table is stored under the prefix 0x00 00 00 00 01. The maximum number of keys held under this configuration is $2^{24} - 1 = 16,777,215$ (1 is subtracted to account for the length value occupying a single space). This is one less than the total number of procedures that can be held by the kernel (which is $2^{24} = 16,777,216$).

0x00 00 00 00 01 + 24 more bytes + 3 zero bytes

Table 2: Procedure table.

Storage Location	Description
0x00 00 00 00 01	length/number of procedures
0x00 00 00 00 01 + keyIndex	procedure key/name: 24 bytes
⋮	capabilities in series
0x00 00 00 00 01 + n	procedure key/name for procedure <i>n</i>

3.4 Procedure Heap


Procedure data is stored using the key 0x00 00 00 00 00. This is combined with the procedure key to produce the following:

0x00 00 00 00 00 + 24 bytes procedure key + 3 bytes data offset

This leaves 3 bytes at the end of the storage location key. If the first of these bytes is 0x00, the storage key refers to various metadata of the procedure, including the location of its contract. If the first byte is any other value, then the storage key corresponds to the list of capabilities of that type. For example, the type of the write procedure is 7, so if the first of the last 3 bytes is 0x07, then that storage key refers to a capability of the write type.

The second value is the index into the capability list. If the second value is 0x00, this indicates the number of capabilities in this list. Therefore the capabilities start at 0x01. For example, if the third last and second last values are 0x0703, this refers to the third write capability (capability at index 0x02) held by this procedure.

The very last value is an offset into the capability data. The meaning of this is different for each capability as they each have different formats. The format of each capability is specified in [Section 4](#)



The following table outlines what data is found at each of the different locations specified by the *Data Offset*.

Table 3: Procedure data.				
Data Offset				Description
0x	00	00	00	address: 20 bytes, aligned right in the 32 bytes
0x	<i>ty</i>	00	00	The number of capabilities of type 0x <i>ty</i> .
0x	<i>ty</i>	<i>in</i>	<i>of</i>	Capability of type 0x <i>ty</i> , with index 0x <i>in</i> – 1, and offset 0x <i>of</i> into that capability.

4 System Calls and Capabilities

All system calls are transactions to the kernel. The transaction data is defined as follows:

Table 4: System call structure.	
Byte Offset	Description
0x0	1 byte, system call value
0x8	The system call data as defined by each system call
⋮	

If the transaction data is longer than specified by the system call format, the additional data is simply ignored. **TODO:** We should consider throwing an error on this.

System call numbers match capability numbers.

Table 5: System call and capability numbers.	
Type Value	Description
0x0	Null capability / noop
0x3	Call Procedure
0x4	Register Procedure
0x5	Delete Procedure
0x6	Set Entry Procedure
0x7	Write
0x8	Log
0x9	Gas Send

4.1 Executing a System Call

The running code will issue a transaction to the kernel by performing a DELEGATECALL to the kernel with the transactions message data as both above and in the definitions of each capability. The system call will be processed as follows:

1. The kernel will receive the transaction as defined by the EVM.
2. The kernel will read the first byte of the transaction data as an unsigned 8-bit integer. This value is s .
3. If s is not one of the values listed in Table 5 then the kernel will revert the current transaction (i.e. the DELEGATECALL) and return the error code `SYSCALL_NOEXIST` (error codes are defined in Section 4.2).
4. If s is 0, it will do nothing and return successfully.
5. Otherwise the control will be passed to the code in the kernel corresponding to the s . For example if the $s = 0x4$, then the code for executing a Register Procedure syscall will be called.
6. This code will execute and return whatever value it deems appropriate.

4.2 Return and Error Codes

This section defines return and error codes for interacting with the kernel. These values are returned by system calls *only* in the event of a failure of the kernel, in which case the DELEGATECALL will have been reverted. As a result, these error codes are provided only when the kernel reverts, if the DELEGATECALL returns normally, whatever is returned is determined by the procedure being called. That is, the DELEGATECALL should leave a 0 on the stack.

The error code is always 8-bits, and may followed by additional data.

Therefore:

- If 1 is left on the stack after the DELEGATECALL of a system call, then the system call succeeded, and whatever data was returned was returned by system call.
- If 0 is left on the stack after the DELEGATECALL of a system call, then the system call failed, and whatever data was returned is defined in the table below.

Table 6: Return and error codes.

Type	Value	Additional Data	Description
SYSCALL_BADCAP	33	None	Capability insufficient
SYSCALL_NOGAS	44	None	Procedure execution ran out of gas.
SYSCALL_REVERT	55	Returned data from procedure	The called procedure reverted.
SYSCALL_FAIL	66	System call specific error data	The system call failed for reasons specific to the system call.
SYSCALL_NOEXIST	111	None	The syscall integer value is not a valid syscall as defined in Table 5

4.3 Capability Subsets

Capability B is a subset¹ of Capability A , if everything that can be done using Capability B can be done using Capability A . The following rules apply:

- A is always a subset of itself. $A \subseteq A$
- Transitivity: $A \subseteq B \wedge B \subseteq C \implies A \subseteq C$

As the nature of each capability is different, the way subsets are defined is different for each one. The following sections will define how subsets are defined for each capability.

4.4 Call System Call

4.4.1 System Call Format

Type Value: 3

Data: Note that here the offset is from after system call type.

Note that here we could pack the data much more densely, but we don't in order to err on the side of simplicity and stick as close as we can to 32-byte values.

¹A subset, not a *strict* or *proper* subset.

Table 7: Call procedure system call format.

Byte Offset	Description
0x0	Capability index; 32 bytes
0x20	Prefix size, 1 bytes
0x21	8 unused bytes
0x28	Procedure key, 24 bytes
End	

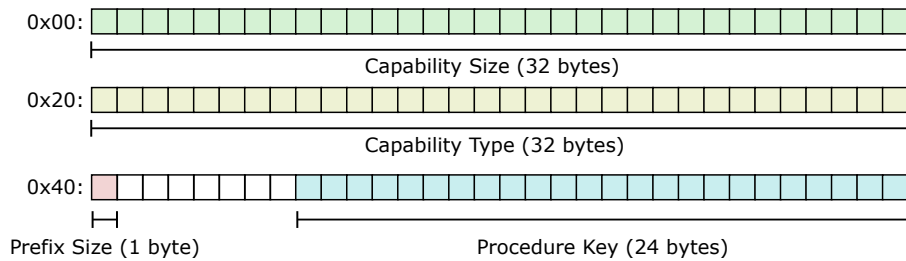
4.4.2 Capability Format

The capability format for the Call Procedure system call defines the a range of procedure keys what the capability allows one to call. This is defined as a base procedure key b and a prefix s . Given this capability, a procedure may call any procedure where the first s bits of the key of that procedure key are the same as the first s bits of procedure key b .

The values of this capability are packed into a single 32-byte value.

Table 8: Call procedure capability format.

Key Offset	Value	Description
0x00	-	The prefix and base procedure key, the first byte is the prefix size which is in the interval $[0,24]$, bytes $[0x1,0x7]$ are unused and undefined, bytes $[0x8,0x32]$ are the 24 bytes of the base procedure key.
End		



4.4.3 Capability Subsets

As defined in Section 4.3, each capability has subsets and supersets. For the call capability a subset is capability which possesses the ability to call no more procedures than the original.

The prefix size of a capability C is given as C_s and the base procedure key C_b . Given a call procedure capability A and a proposed subset B :

- If the prefix size of B is equal to or greater than A (that is: $B_s \geq A_s$), and
- If the first B_s bits of B_b is equal to the first A_s bits of A_b .

4.4.4 Error Codes

If the Call Procedure system call fails for any reason other than those general system call failure conditions, the system call will return **SYSCALL_FAIL** (as defined in Section 4.2) followed by one of the following codes.

Table 9: Return and error codes.

Type	Value	Additional Data	Description
CALL_NOPROC	33	None	The procedure key specified does not exit.

Non-Normative

For example, if the Call Procedure system call was executed with a procedure key that does not exist then the DELEGATECALL will leave a 0 on the stack and return two bytes (0x6633).

4.5 Register Procedure System Call

4.5.1 System Call Format

Type Value: 4

Data: min(96 bytes) max(96 bytes)

Table 10: Register procedure system call structure.

Byte Offset	Description
0x0	Capability index; 32 bytes
0x20	Procedure key; 24 bytes in 32 bytes
0x40	Procedure address; 20 bytes in 32 bytes
⋮	Capabilities; A series of 32 byte values which correspond exactly to the keys as defined in Table 11. The length of this list is not provided.
End	

Table 11: Capability data as sent in the Register Procedure system call.

Key Offset	Description
0x00	CapSize: 2 bytes, aligned right in the 32 bytes, the given name
0x01	CapType: 1 byte, aligned right in the 32 bytes
0x02	key value #1: 32 bytes
⋮	capabilities in series
CapSize	final key value

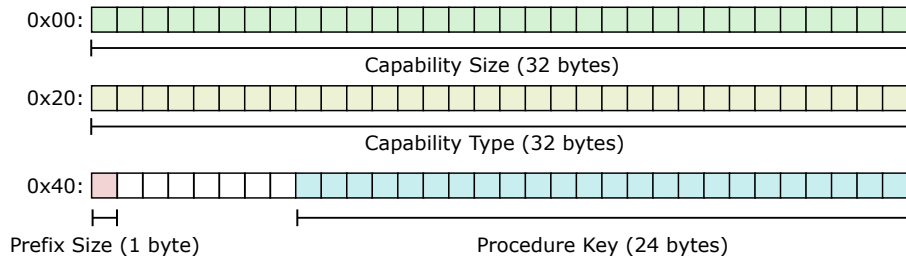
4.5.2 Capability Format

The capability format for the Register Procedure system call defines the a range of procedure keys what the capability allows one to call. This is defined as a base procedure key b and a prefix s . Given this capability, a procedure may call any procedure where the first s bits of the key of that procedure key are the same as the first s bits of procedure key b .

NB: Register procedure also relies on all of the other capabilities possessed by a procedure to determine what capabilities the new procedure will have.

Table 12: Register procedure capability format.

Key Offset	Value	Description
0x00	-	The prefix and base procedure key, the first byte is the prefix size which is in the interval $[0,24]$, bytes $[0x1,0x7]$ are unused and undefined, bytes $[0x8,0x32]$ are the 24 bytes of the base procedure key.
End		



4.5.3 Capability Subsets

As defined in Section 4.3, each capability has subsets and supersets. For the register procedure capability a subset is capability which possesses the ability to

register no more procedures than the original.

The prefix size of a capability C is given as C_s and the base procedure key C_b . Given a register procedure capability A and a proposed subset B :

- If the prefix size of B is equal to or greater than A (that is: $B_s \geq A_s$), and
- If the first B_s bits of B_b is equal to the first A_s bits of A_b .

4.5.4 Function

This registers a new procedure which already exists as a contract at a certain address, giving it a name and a list of capabilities. It does a number of things:

- Validate the code at the given address to show that it complies with the requirements of procedure code as described in Section 5.
- Add the procedure name to the procedure list.
- Store the specified capabilities with the procedure, but only if each of the capabilities is found to be a subset of one of the capabilities of the procedure performing this system call.

There are two highly critical functions here upon which the safety of the kernel relies. The first is the validation of the procedure bytecode, which is covered in Section 5, but it is also necessary to ensure that the capabilities that are being asked to be given to the new procedure can be provided by the current procedure. In order to satisfy this constraint it must be shown that for every requested procedure, there exists a capability in the capability list of the current procedure. It is currently not able to “combine” capabilities. That is, if the procedure has a capability of `Write(0x80,5)` and `Write(0x85,5)` it cannot provide a capability `Write(0x80,10)` even though it theoretically still able to perform the same writes.

The algorithm is as follows:

1. For each of the capabilities in the list of requested capabilities, search through the list of the current procedures capabilities until a superset of the requested capability is found.
2. If for any of the requested capabilities a superset is not found, abort the entire process.
3. If for all of the requested capabilities a superset is found, register the capability.

4.5.5 Error Codes

If the Register Procedure system call fails for any reason other than those general system call failure conditions, the system call will return `SYSCALL_FAIL` (as defined in Section 4.2) followed by one of the following codes.

Table 13: Return and error codes.

Type	Value	Additional Data	Description
REG_TOOMANYCAPS	77	None	Too many caps were provided.

Non-Normative

For example, if the Register Procedure system call was executed but provides too many capabilities then the `DELEGATECALL` will leave a 0 on the stack and return two bytes (0x6677).

4.6 Delete Procedure System Call

4.6.1 System Call Format

Type Value: 5

Data: min(64 bytes) max(64 bytes)

Table 14: Delete procedure system call.


Byte Offset	Description
0x0	Capability index; 32 bytes
0x20	Procedure key; 24 bytes in 32 bytes
End	

4.6.2 Capability Format

The capability format for the Delete Procedure system call defines the a range of procedure keys what the capability allows one to delete. This is defined as a base procedure key b and a prefix s . Given this capability, a procedure may delete any procedure where the first s bits of the key of that procedure key are the same as the first s bits of procedure key b .

Table 15: Delete procedure capability format.

Key Offset	Value	Description
0x00	-	The prefix and base procedure key, the first byte is the prefix size which is in the interval $[0,24]$, bytes $[0x1,0x7]$ are unused and undefined, bytes $[0x8,0x32]$ are the 24 bytes of the base procedure key.
End		



4.6.3 Capability Subsets

As defined in Section 4.3, each capability has subsets and supersets. For the delete procedure capability a subset is capability which possesses the ability to delete no more procedures than the original.

The prefix size of a capability C is given as C_s and the base procedure key C_b . Given a delete procedure capability A and a proposed subset B :

- If the prefix size of B is equal to or greater than A (that is: $B_s \geq A_s$), and
- If the first B_s bits of B_b is equal to the first A_s bits of A_b .

4.7 Set Entry Procedure System Call

The value of the Entry Procedure is 24-byte procedure key which identified which procedure should be first called upon receiving a transaction. This is currently store at the storage location:

0x00 00 00 00 04 00 ... 00 00 00 00

4.7.1 Error Codes

If the Delete Procedure system call fails for any reason other than those general system call failure conditions, the system call will return **SYSCALL_FAIL** (as defined in Section 4.2) followed by one of the following codes.

Table 16: Return and error codes.

Type	Value	Additional Data	Description
DEL_NOPROC	33	None	The procedure key specified does not exist.

Non-Normative

For example, if the Delete Procedure system call was executed with a procedure key that does not exist then the DELEGATECALL will leave a 0 on the stack and return two bytes (0x6633).

4.7.2 System Call Format

Type Value: 6

Data: min(64 bytes) max(64 bytes)

Table 17: Set entry procedure system call format.

Byte Offset	Description
0x0	Capability index; 32 bytes
0x20	Procedure key; 24 bytes in 32 bytes
End	

4.7.3 Capability Format

This capability has no associated data and is therefore of size zero.

4.7.4 Capability Subsets

As defined in Section 4.3, each capability has subsets and supersets. As the Set Entry Procedure capability is a trivial value, there is only one value for Set Entry Procedure capabilities and they are all equal.

Given a Set Entry Procedure capability A and any other Set Entry Procedure capability B , not only is $B \subseteq A$ but also $B = A$.

4.7.5 Function

The only possible state change is the "entryProcedure" value. No other state changes should occur.

4.7.6 Error Codes

If the Set Entry Procedure system call fails for any reason other than those general system call failure conditions, the system call will return `SYSCALL_FAIL` (as defined in Section 4.2) followed by one of the following codes.

Table 18: Return and error codes.

Type	Value	Additional Data	Description
SETENT_NOPROC	33	None	The procedure key specified does not exit.

Non-Normative

For example, if the Set Entry Procedure system call was executed with a procedure key that does not exist then the `DELEGATECALL` will leave a 0 on the stack and return two bytes (0x6633).

4.8 Write System Call

This system call will write a single 32-byte value under a single 32-byte key in the storage of the kernel instance.

4.8.1 System Call Format

Type Value: 7

Data:

Table 19: System call structure.

Byte Offset	Description
0x0	Capability index; 32 bytes
0x20	Write address; 32 bytes
0x40	Write values; 32 bytes
End	

4.8.2 Capability Format

The write capability includes 2 values: the first is the base address where we can write to storage. The second is the number of *additional* addresses we can write to. For example, if the first value is 0x8000, and the second value is 0, we can write only to location 0x8000. If the second value was 5, we could write to 0x8000, 0x8000+1, 0x8000+2, 0x8000+3, 0x8000+4, and 0x8000+5.

Table 20: Write capability format.

Key Offset	Value	Description
0x00	a	The base storage address
0x01	n	The number of additional keys we can write to
End		

4.8.3 Capability Subsets

As defined in Section 4.3, each capability has subsets and supersets.

Given a Write capability A and a proposed subset B , where C_a is the base address of capability C and C_n is the number of additional keys writable with capability C , $B \subseteq A \iff (B_a \geq A_a) \wedge (B_a + B_n \leq A_a + A_n)$.

Or more verbosely, B is only a subset of A if and only if:

- The lowest writable address (which is the base address) of B is greater than or equal to the lowest writable address of A , and
- The highest writable address (which is base address plus the number of additional keys) of B is less than or equal to the highest writable address of A .

4.8.4 Error Codes

There are no specific error conditions for SSTORE, therefore only general system call errors are applicable.

4.9 Log System Call

4.9.1 System Call Format

Type Value: 8

Data: min(96 bytes) max(224 bytes)

Table 21: System call structure.

Byte Offset	Description
0x0	Capability index; 32 bytes
0x20	Number of topics (nTopics)
...	Potentially topic value #1; 32 bytes
...	Potentially topic value #2; 32 bytes
...	Potentially topic value #3; 32 bytes
...	Potentially topic value #4; 32 bytes
0x40+(nTopics x 0x20)	Log value 32 bytes
End	

4.9.2 Capability Format

The Write capability includes between 0 and 4 values. Each value forces the use of a particular value. For example: if the capability has 1 value, then that means that the first topic in the log call must equal that. If it has 2 values then the first log topic must be the first of those values and the second topic must be the second of those values and so on. If there are no topics listed then there are no restrictions on what can be logged.

This capability is 5 values. If the number of enforced topics is less than 4, then the unused values are undefined.

Table 22: Log capability format.

Key Offset	Value	Description
0x00	[0x00, 0x4]	Number of enforced topics; 32 bytes
0x01	t_1	Potentially an enforced value for the first topic
0x02	t_2	Potentially an enforced value for the second topic
0x03	t_3	Potentially an enforced value for the third topic
0x04	t_4	Potentially an enforced value for the fourth topic
End		

4.9.3 Capability Subsets

As defined in Section 4.3, each capability has subsets and supersets.

Given a Log capability A and a proposed subset B , $B \subseteq A$ if and only if for each log topic of A the topic is either undefined or equal to that of B .

If C_{t_i} , where i is 1, 2, 3, or 4, is the log topic i of capability C , and C_n is the number of defined topics:

$$\forall i \in \{1, 2, \dots, A_n\} . B_{t_i} = A_{t_i}$$

4.9.4 Error Codes

There are no specific error conditions for SSTORE, therefore only general system call errors are applicable.

4.10 Value Send System Call

This system call sends value (Ether) to another account.

4.10.1 System Call Format

Type Value: 9

Data:

Table 23: System call structure.

Byte Offset	Description
0x0	Capability index; 32 bytes
0x20	Account address; 20-bytes right aligned in 32 bytes
0x40	Value amount; 32 bytes
End	

4.10.2 Capability Format

This value contains no data other than its presence, therefore has no format.

4.10.3 Capability Subsets

As defined in Section 4.3, each capability has subsets and supersets. As the Value Send capability is a trivial value, there is only one value for Value Send capabilities and they are all equal.

Given a Value Send capability A and any other Value Send capability B , not only is $B \subseteq A$ but also $B = A$.

5 Procedure Bytecode Validation

In order to pass bytecode validation each opcode in the contract must either be on the whitelist, or part of a recognised system call. The whitelist comprises all non-state-changing opcodes, and a system call is in the form:

Listing 1: Sequence of steps to perform a system call.

```
CALLER          // Get Caller
GAS             // Put all the available gas on the stack
DELEGATECALL    // Delegate Call to Caller
```